

# Economic AI

---

Greg Lewis (MSR + NBER)   Matt Taddy (MSR + Chicago)  
Jason Hartford (UBC)   Kevin Leyton-Brown (UBC)  
Kui Tang (Columbia)   Dave Blei (Columbia)  
Matt Goldman (MSFT)   Justin Rao (MSR)   Di Wang (MSFT)  
James Zou (Stanford)   Mengting Wan (UCSD)   Richard Li (UW)

# What do economists do?


320

JUNE 1980

THE AMERICAN ECONOMIC REVIEW  
TABLE 2—TOTAL EXPENDITURE AND OWN-PRICE ELASTICITIES


	Levels Model		First-Differences Model	
	Unconstrained $e_i$	Homogeneous $e_{ii}$	Unconstrained $e_i$	Homogeneous $e_{ii}$
Food	0.21	0.07	0.04	0.17
Clothing	2.00	-0.92	2.83	2.92
Housing	0.30	-0.31	0.04	-0.94
Fuel	1.67	-0.28	1.00	-0.31
Drink and Tobacco	1.22	-0.60	1.37	0.00
Transport and Communication	1.23	-1.21	1.14	-1.23
Other goods	1.21	-0.72	2.03	-0.52
Other services	1.40	-0.93	1.03	-0.78

GOOD 2




A

GOOD 2



B

GOOD 2



C

**Income expansion paths.** Panel A depicts unit elastic demands, in panel B good 2 is a luxury good, and in panel C, good 1 is an inferior good.

# What do they need to do today?

The image is a composite of two screenshots. The background is a Google search results page for the query "toddler shoes". The search bar at the top contains the text "toddler shoes" and a magnifying glass icon. Below the search bar, there are tabs for "Web", "Images", "Videos", "Maps", "News", and "Explore". The "Web" tab is selected. Below the tabs, it says "91,800,000 RESULTS" and "Any time". The first search result is "Shop DSW Kids Shoes | dsw.com" with the URL "www.dsw.com/kids" and the text "DSW, Inc.". Below this, there is a description: "The Latest Kids Styles @ Participating DSW Stores Today! ... and best brands for infants, toddlers, and ...". There are also links for "Sign Up for DSW® Rewards" and "Find a Store Near You".

Overlaid on the bottom left of the Google search page is a screenshot of the Microsoft Azure portal. The window title is "Microsoft Azure" and the subtitle is "New HDInsight Cluster". The main content area is titled "Cluster Type configuration" and includes a "Report a bug" button. Below this, there is a section for "Cluster Type" with a dropdown menu set to "Spark". To the right, there is a section for "Operating System" with a dropdown menu set to "Linux". Below these, there are two columns of configuration options. The left column is for the "STANDARD" tier and includes "Administration" (Manage, monitor, connect), "Scalability" (On-demand node scaling), and "99.9% Uptime SLA". The right column is for the "PREMIUM (PREVIEW)" tier and includes "Administration" (Manage, monitor, connect) and "Scalability" (On-demand node scaling). At the bottom of the Azure window, there is a "Version" dropdown menu set to "Spark 2.0.0 (HDI 3.5)".

# What can we ML do to help?

The **dimension and complexity** of the problem space has exploded

We can develop ML to navigate this space: **stay safe and automate**

# Example: Demand System

Suppose that you have transactions ' $t$ ' on products ' $j$ '.

Write the quantity bought ' $q$ ' as

$$q_{tj} = \alpha_{tj} + \gamma_j p_{tj} + e_{tj}$$

a function of utility we can ( $\alpha_{tj}$ ) and can't ( $e_{tj}$ ) see, plus price  $p_{tj}$ .

You need to have a model like this to target customers or set prices.

# But it's a system!

For example: There many different products

Demand for  $j$  depends on **substitutes** and **complements**

Or: where does price come from?

$$\log p_{tj} = \varphi_{tj} + \psi_j q_{tj}^* + v_{tj}$$

and the *demand system* is in equilibrium when  $q_{tj}^* = q_{tj}$

# Product Co-occurrence

Ignoring price and observable demand shifters...

$$q_{tj} = \alpha_{tj} + \gamma_j \log p_{tj} + \varepsilon_{tj}$$

Classic “data mining” seeks *association rules* in *market baskets*:

Find  $j \neq k$  pairs so that  $\mathbb{E}[q_{tj}q_{tk}] \gg \mathbb{E}q_{tj}\mathbb{E}q_{tk}$

We can do better

# From word to product embedding

Words are like products and sentences are like baskets

Tools like **word2vec** and **glove** map from the  $J$ -dimensional discrete space to [two] vector *embeddings* in  $\mathbb{R}^S$ , and  $S \ll J$

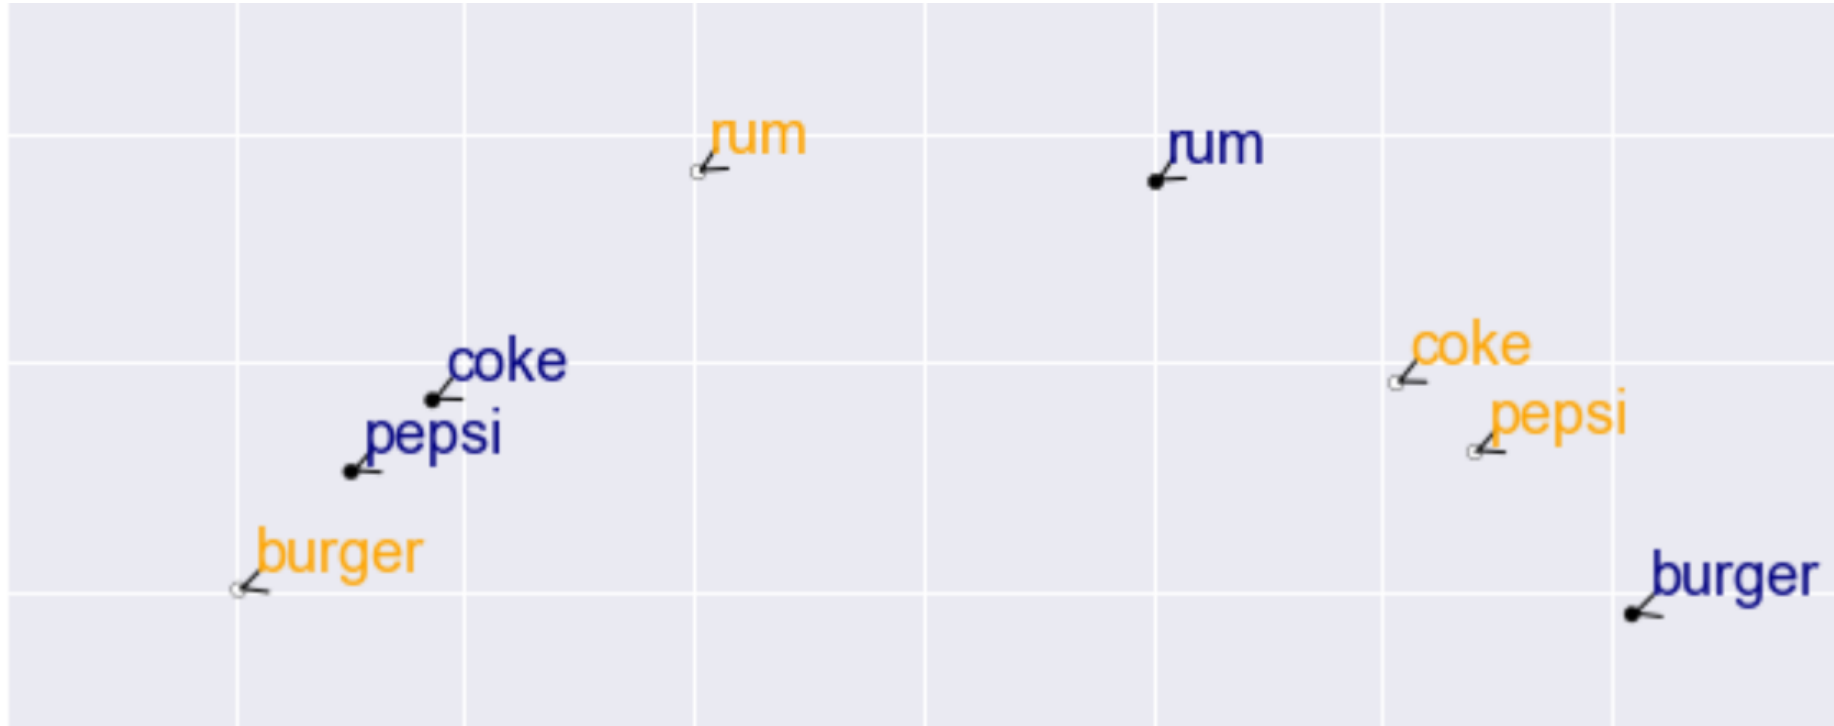
$$\max \left\{ \sum_t \sum_{j,k} q_{tj} q_{tk} \mathbf{u}'_j \mathbf{v}_k - A(\mathbf{U}, \mathbf{V}, \mathbf{C}) \right\}$$

where  $c_{jk} = \sum_t q_{tj} q_{tk}$  and  $A(\cdot)$  is a normalizing constant

W2V uses logit model, Glove minimizes  $\sum_{j,k} w_{jk} (c_{jk} - \mathbf{u}'_j \mathbf{v}_k)^2$ ,  $w_{jk} = \mathbb{1}[c_{jk} > 0]$



# Product Embeddings



substitutes (synonyms) are close in the same vector space  
complements (topical words) are close across vector spaces

# Moving to a demand system (AID)

It's *almost* ideal:

$$\mathbf{s}_t = \boldsymbol{\alpha} + \boldsymbol{\Gamma} \log(\mathbf{p}_t) + \boldsymbol{\beta} \log \frac{e_t}{\phi_t} + \boldsymbol{\varepsilon}_t$$

$s_{tj}$  is the **budget share** for product  $j$  in basket  $t$  and  $e_t$  is the budget  
( $e_t = \sum_j \$_{tj}$  and  $s_{tj} = \$_{tj}/e_t$ )

$\phi_t$  is the **translog price index**  $\sum_j \log p_{tj} [\alpha_j + \sum_k \gamma_{jk} \log p_{tk}]$

This is meaningful after aggregation, and it's *easier* to estimate...

# But not so easy...

In grocery stores, we have 100,000s of products

Each week, at each store, for each product,  
we will have different demand and elasticity and income effects.

$$\mathbb{E} s_{tj} = \alpha_{tj} + \sum_k \gamma_{tjk} \log p_{tj} + \beta_{tj} \frac{e_t}{\phi_t}$$

# Featurize and Embed

Grab massive  $\mathbf{x}_{tj}$  that *hierarchically* encode products and transactions

$$\mathbb{E}s_{tj} = \mathbf{x}'_{tj}\boldsymbol{\alpha} + \mathbf{x}'_{tj}\boldsymbol{\delta} \log p_{tj} + \mathbf{x}'_{tj}\boldsymbol{\beta} \log \left( \frac{e_t}{\phi_t(\mathbf{x}_{tj})} \right) + \boldsymbol{\Gamma} \log \mathbf{p}_t$$

and use a symmetric square matrix factorization  $\boldsymbol{\Gamma} = \mathbf{U}\mathbf{V}' + \mathbf{V}\mathbf{U}'$


*Note that we are using the actual translog and solving a bilinear system*

# Beer Store own-price (compensated) elasticities


*one shared:  $x_{tj} = 1$*

$$\frac{dq}{dp} \frac{p}{q} = -0.23$$

*brand-specific:  $x_{tjk} = \mathbb{1}_{[k=j]}$*




*$x_{tj}$  = featurized description*



# Beer Store cross-product compensated elasticities

*Elasticity matrix (omitting diagonal)*



But wait... it's still a system

$$s_t = \alpha + \Gamma \log(\mathbf{p}_t) + \beta \log \frac{e_t}{\phi_t} + \mathbf{e}_t$$

Recall: where does price come from?

$$\log p_{tj} = \varphi_{tj} + \psi_j q_{tj}^* + v_{tj}$$

and the *demand system* is in equilibrium when  $q_{tj}^* = q_{tj}$

This equilibrium introduces 'price endogeneity':  $\mathbb{E}[p_{tj} e_{tj}] \neq 0$

# Endogenous Errors

$$y = g(p, \mathbf{x}) + e \text{ and } \mathbb{E}[p e] \neq 0$$


If you estimate this using naïve ML, you'll get

$$E[y|p, \mathbf{x}] = E_{e|p}[g(p, \mathbf{x}) + e] = g(p, \mathbf{x}) + E[e|t, \mathbf{x}]$$

This works for **prediction**. It doesn't work for **counterfactual** inference:


*What happens if I change  $p$  independent of  $e$ ?*





— Prediction — Counterfactual Prediction


# Instrumental Variables (IV)



In IV we have a special  $z \perp e$  that influences policy  $p$  but not response  $y$ .

- Supplier costs that move price independent of demand (e.g., fish, oil)
- Any source of treatment randomization (intent to treat, AB tests, lottery)

# Instrumental Variables (IV)



The *exclusion structure* implies

$$E[y|x, z] = E[g(p, x) + e|x, z] = \int g(p, x) dF(p|x, z)$$

So to solve for *structural*  $g(p, x)$  we have a new learning problem

$$\min_{g \in G} \sum \left( y_i - \int g(p, x_i) dF(p|x_i, z_i) \right)^2$$

$$\min_{g \in G} \sum \left( y_i - \int g(p, x_i) dP(p|x_i, z_i) \right)^2$$

2SLS:

$$p = \beta z + v \text{ and } g(p) = \tau p \text{ so that } \int g(p) dP(p|z) = \tau \hat{p} = \tau \hat{\beta} z$$

So you first regress  $p$  on  $z$  then regress  $y$  on  $\hat{p}$  to recover  $\hat{\tau}$ .

This requires strict assumptions and homogeneous treatment effects.

$$\min_{g \in G} \sum \left( y_i - \int g(p, x_i) dP(p|x_i, z_i) \right)^2$$

Or look to nonparametric 2SLS like in Newey and Powell:

$$g(p, x_i) \approx \sum_k \varphi_k(p, x_i) \text{ and } \varphi_k(p, x_i) \approx \sum_j \phi_{kj}(x_i, z_i)$$

But this requires careful crafting and will not scale with  $\dim(x)$

$$\min_{g \in G} \sum \left( y_i - \int g(p, x_i) dF(p|x_i, z_i) \right)^2$$


Instead, we propose to **target the integral loss function directly**

For discrete treatment

- Fit distributions  $\hat{F}(p|x_i, z_i)$  with probability masses  $\hat{f}(p_b|x_i, z_i)$
- Train  $\hat{g}$  to minimize  $\left[ y_i - \sum_b g(\hat{p}_b, x_i) \hat{f}(p_b|x_i, z_i) \right]^2$

And you've turned IV into two *generic* machine learning tasks

# Learning to love Deep Nets




# What is a deep net?

$$\hat{y}_i = \sum_k h_k^0(a_{ik}), \quad a_{ik} = \sum_j w_{kj} z_j, \quad z_j = \sum_l h_l^1(b_{il}), \dots$$

And so-on until you get down to a bottom layer  $\{f_l(x_i)\}_l$

Many different variations here: recursive, convolutional, ...



Apart from the bottom, usually  $h(v) = \max\{0, v\}$





e.g., first-stage learning for  $F(p|x_i, z_i)$

Bishop 96: Final layer of network parametrizes a mixture of Gaussians



The second stage involves an integral loss function

If  $p$  is not discrete or can take many values, we can't just integrate

Brute force just samples from  $\hat{F}(p|x_i, z_i)$  and you take gradients on

$$\frac{1}{N} \sum_i \left( y_i - \frac{1}{B} \sum_b g(\hat{p}_{ib}, x_i; \theta) \right)^2, \quad \hat{p}_{ib} \sim \hat{F}(p|x_i, z_i)$$

But this is inefficient

And more generally, it's inefficient even *without* the integral...

# Stochastic Gradient Descent

You have loss  $L(\mathbf{D}, \theta)$  where  $\mathbf{D} = [\mathbf{d}_1 \dots \mathbf{d}_N]$

In the usual GD, you iteratively descend

$$\theta_t = \theta_{t-1} - \mathbf{C}_t \nabla L(\mathbf{D}, \theta_{t-1})$$

In SGD, you instead follow *noisy* but *unbiased* sample gradients

$$\theta_t = \theta_{t-1} - \mathbf{C}_t \nabla L(\{\mathbf{d}_{t_b}\}_{b=1}^B, \theta_{t-1})$$

# Why SGD? You get what you need faster

6 Léon Bottou

**Table 2.** Asymptotic equivalents for various optimization algorithms: gradient descent (GD, eq. 2), second order gradient descent (2GD, eq. 3), stochastic gradient descent (SGD, eq. 4), and second order stochastic gradient descent (2SGD, eq. 5). Although they are the worst optimization algorithms, SGD and 2SGD achieve the fastest convergence speed on the expected risk. They differ only by constant factors not shown in this table, such as condition numbers and weight vector dimension.

	GD	2GD	SGD	2SGD
Time per iteration :	$n$	$n$	1	1
Iterations to accuracy $\rho$ :	$\log \frac{1}{\rho}$	$\log \log \frac{1}{\rho}$	$\frac{1}{\rho}$	$\frac{1}{\rho}$
Time to accuracy $\rho$ :	$n \log \frac{1}{\rho}$	$n \log \log \frac{1}{\rho}$	$\frac{1}{\rho}$	$\frac{1}{\rho}$
Time to excess error $\varepsilon$ :	$\frac{1}{\varepsilon^{1/\alpha}} \log \frac{1}{\varepsilon}$	$\frac{1}{\varepsilon^{1/\alpha}} \log \frac{1}{\varepsilon} \log \log \frac{1}{\varepsilon}$	$\frac{1}{\varepsilon}$	$\frac{1}{\varepsilon}$

$$\mathcal{E} = \mathcal{E}_{\text{app}} + \mathcal{E}_{\text{est}} + \mathcal{E}_{\text{opt}} \sim \mathcal{E}_{\text{app}} + \left( \frac{\log n}{n} \right)^\alpha + \rho$$

# Learning rates

$$\theta_t = \theta_{t-1} - \mathbf{C}_t \nabla L(\{\mathbf{d}_{t_b}\}_{b=1}^B, \theta_{t-1}) = \theta_{t-1} - \mathbf{C}_t \nabla L_t$$

What is  $\mathbf{C}_t$ ? Hopefully not  $c_t \mathbf{I}$  ...

Ideally,  $\mathbf{C}_t \rightarrow t^{-1} \mathbf{C}$  and  $\mathbf{C}^{-1} = -\nabla \nabla L(\mathbf{D}, \theta^*)$ , but that's infeasible

*ADAGRAD* (for convex loss)

$$\mathbf{C}_t = \text{diag} \left( \sqrt{\sum_{j=1}^n \nabla L_{tj}^2} \right) \Rightarrow \sum_{t=1}^T (\text{err}(\theta_t) - \text{err}(\theta^*)) = O(\|\theta^*\| \text{tr}(\mathbf{C}_T))$$

We use heuristic innovations on this (ADAM, momentum, etc)

# SGD for integral loss functions

Our one-observation stochastic gradient is

$$\nabla L(d_i, \theta) = -2 \left( y_i - \int g_\theta(p, x_i) d\hat{F}(p|x_i, z_i) \right) \int g'_\theta(p, x_i) d\hat{F}(p|x_i, z_i)$$

Do SGD by pairing each observation with *two independent* treatment draws

$$\nabla \hat{L}(d_i, \theta) = -2(y_i - g_\theta(\hat{p}_{i1}, x_i)) g'_\theta(\hat{p}_{i2}, x_i), \quad \hat{p}_{ib} \sim \hat{F}(p|x_i, z_i)$$

So long as the draws are independent,  $\mathbb{E} \nabla \hat{L}(d_i, \theta) = \mathbb{E} \nabla L(d_i, \theta) = \nabla L(\mathbf{D}, \theta)$

# Lower variance grad $\Rightarrow$ faster/better convergence

e.g., in most architectures you want to fit  $F$  and  $g$  in two-stages

And with MC integration, consider:

$$\nabla \hat{L}(\{\mathbf{d}_b\}_{b=1}^B, \theta) = \sum_{b=1}^B (y_i - g_\theta(\hat{p}_{b1}, x_i)) g'_\theta(\hat{p}_{b2}, x_i)$$

$$\nabla \tilde{L}^B(\mathbf{d}, \theta) = (y_i - \sum_{b=1}^B g_\theta(\hat{p}_{b1}, x_i)) \sum_{b=1}^B g'_\theta(\hat{p}_{b2}, x_i)$$

Both involve the same number of operations, but

$$\text{var} \nabla \hat{L}(\{\mathbf{d}_b\}_{b=1}^B, \theta) = \frac{\text{var} \nabla \hat{L}(\mathbf{d}, \theta)}{B} \quad \text{while} \quad \text{var} \nabla \tilde{L}^B(\mathbf{d}, \theta) \approx \frac{E[\text{var} \nabla \tilde{L}^1(\mathbf{d}, \theta)]}{B} + \text{var}(\nabla L(\mathbf{d}, \theta))$$

Aside: we can use SGD more in econ ...

There are a ton of setups where we use simulation to solve

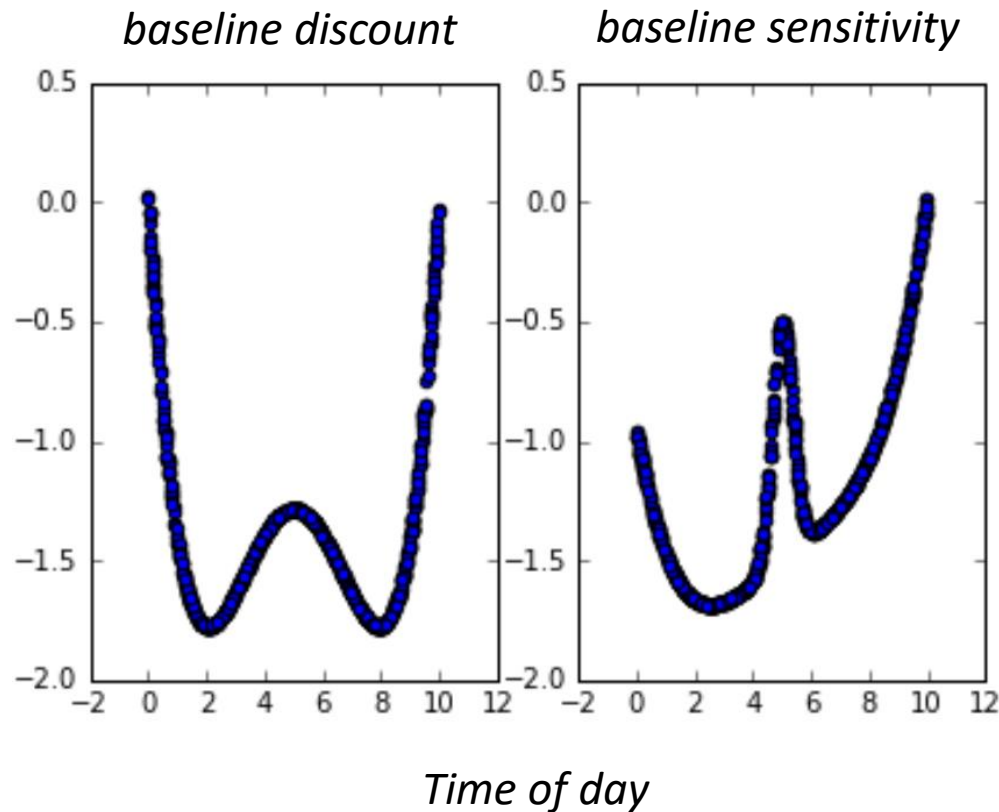
$$\min_{\beta} \sum \left( y_i - \int g(x_i; \theta) dP(\theta | \beta) \right)^2$$

Random coefficient models, or wider class in Pakes and Pollard


Sampling SGD is a perfect fit here



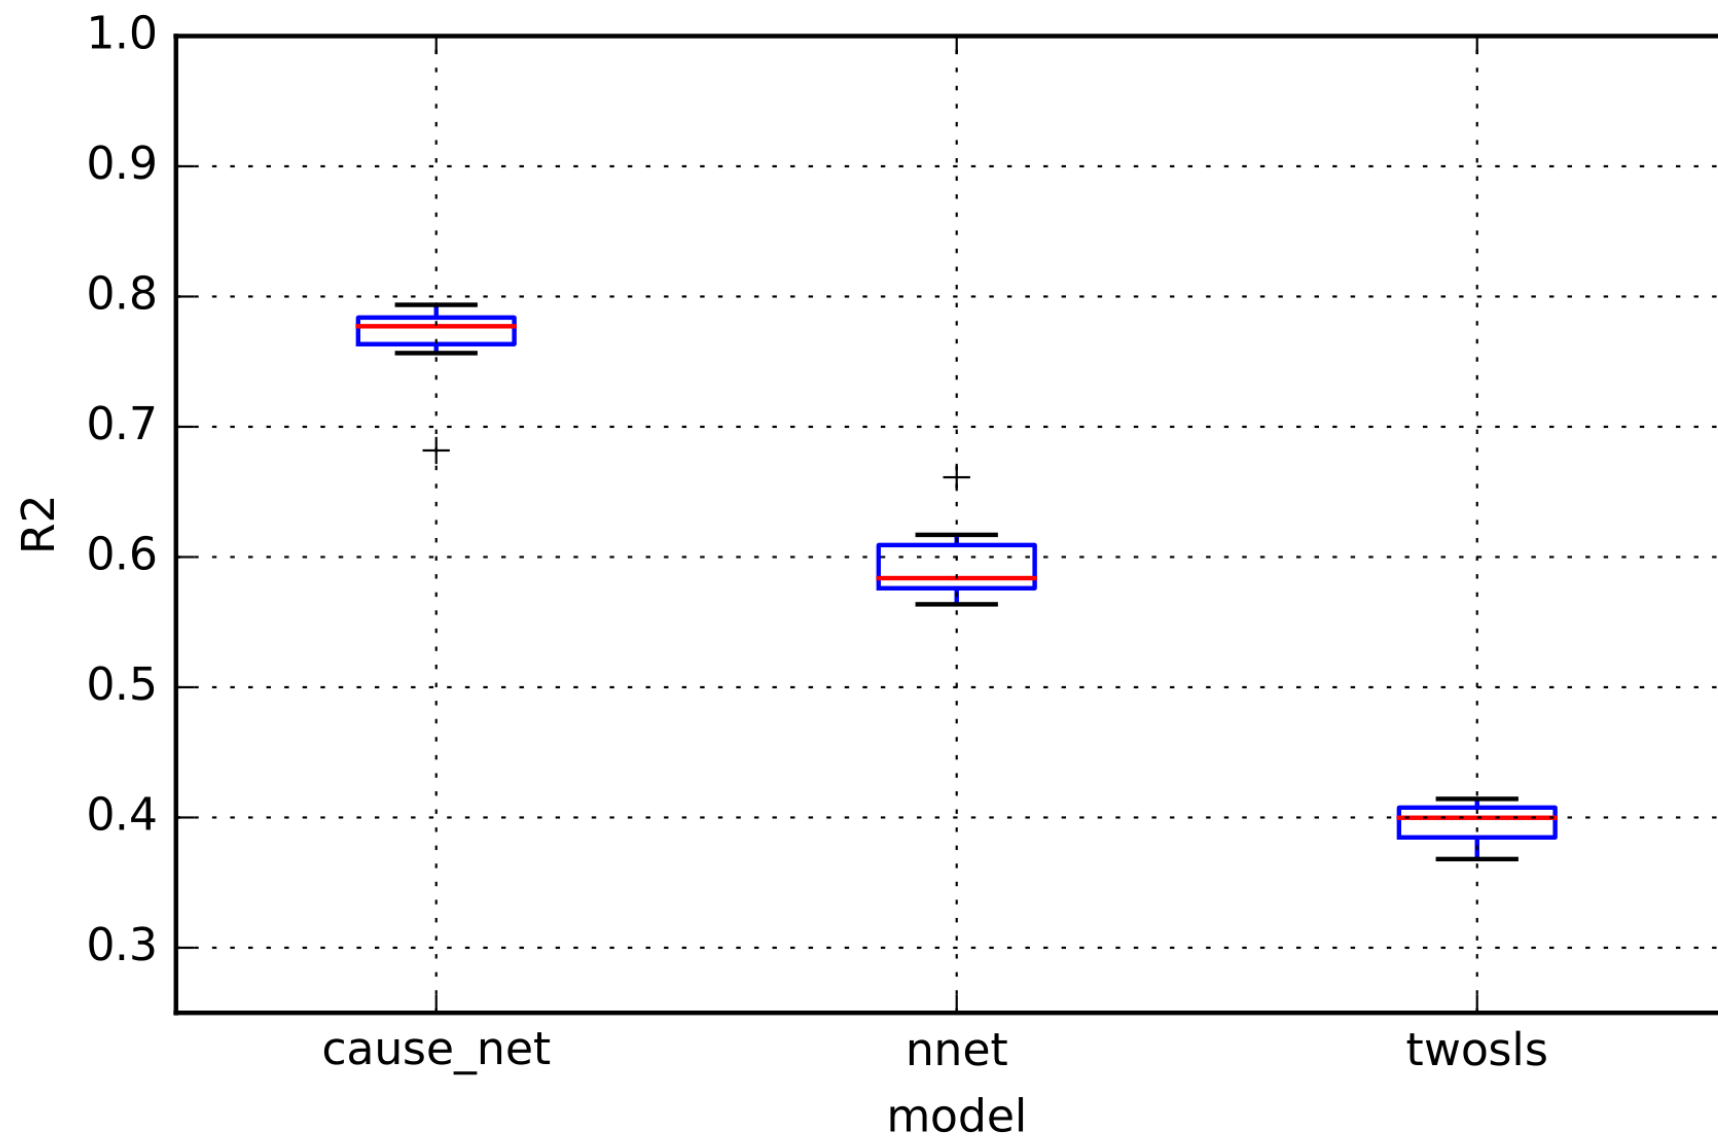
# heterogeneous price effects



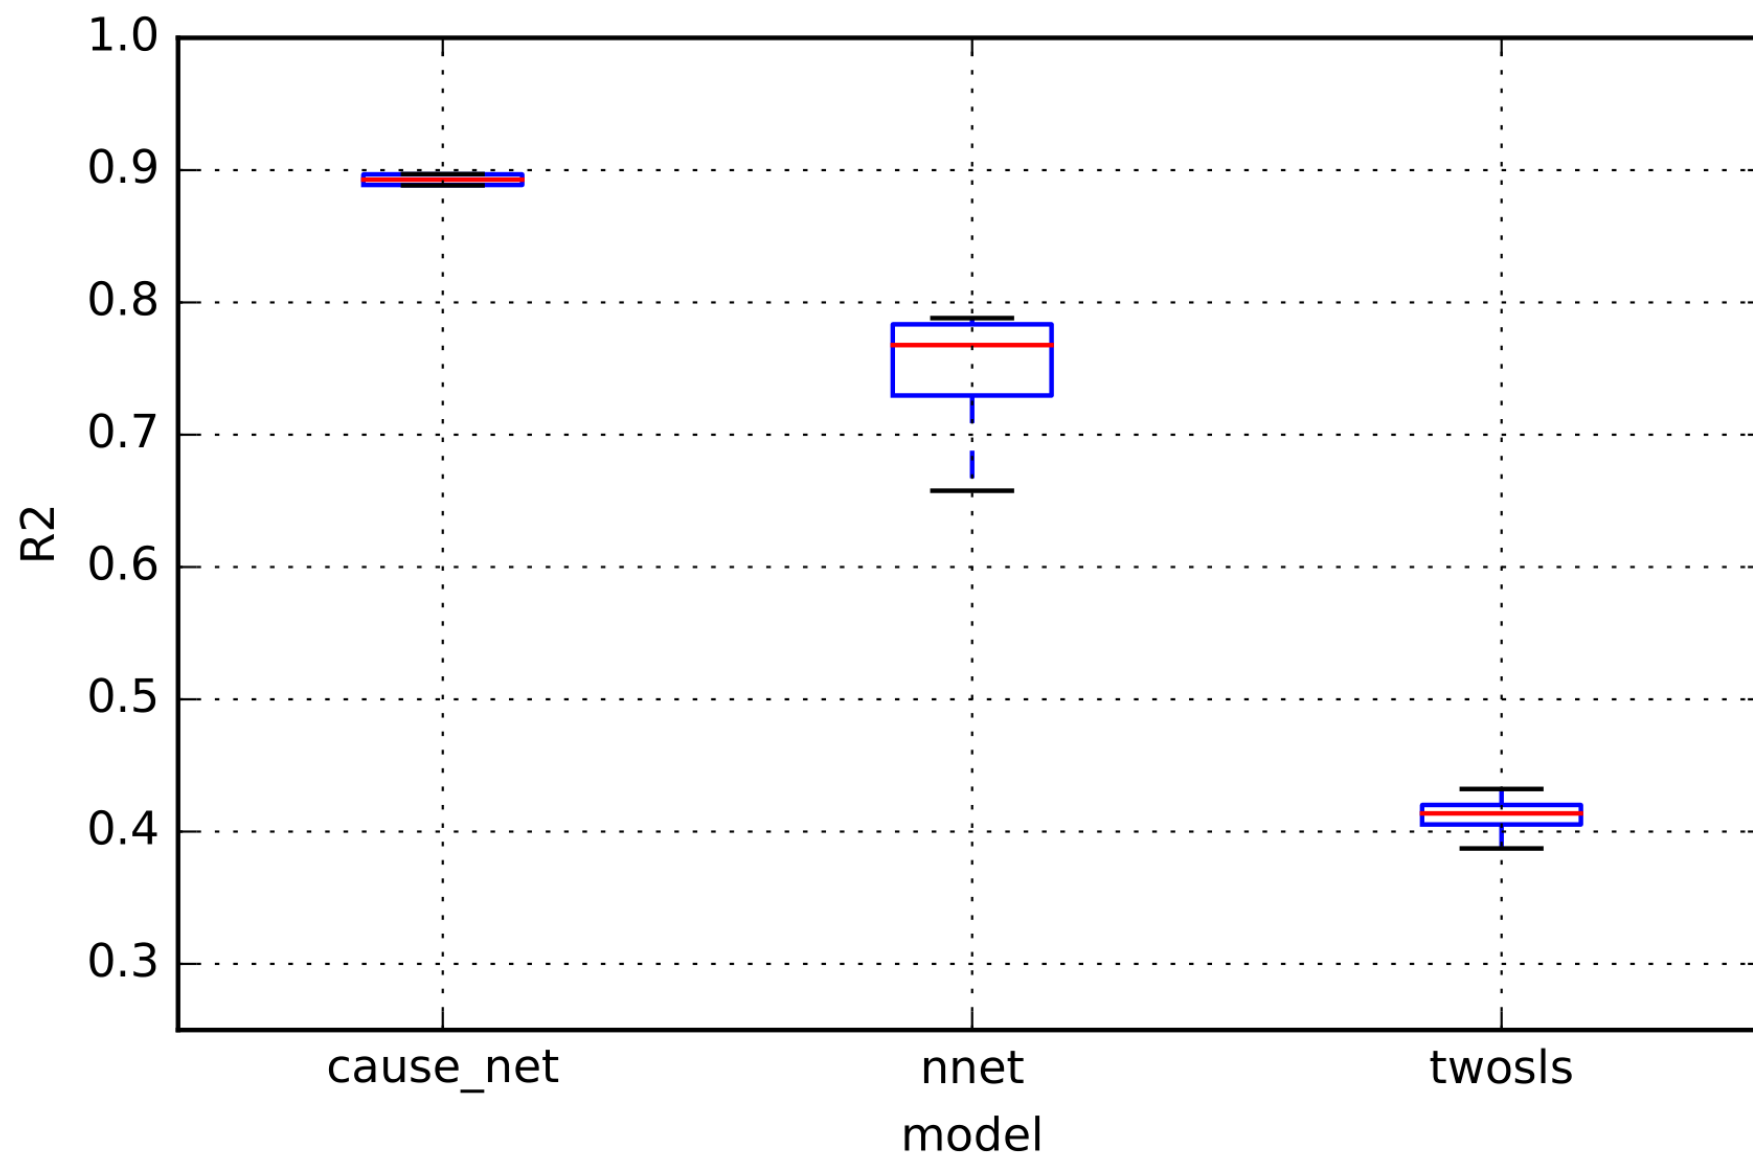
- 'time' dependent prices and elasticity
- 7 different customer types multiply discounts, sensitivities, and demand



N=1000



N=5000



# Identification

The truth  $g_0$  is identified if this equation has a unique solution. Let  $T[g](x, z) = \int g(p, x) dF(p|x, z)$ . Since  $T$  is a linear operator on the function space  $G$ , uniqueness of the solution is equivalent to the kernel of  $T$  being the singleton zero function:

$$\{g \in G : T(g) = 0\} = \{0\}$$

This can equivalently be stated as the “completeness condition” of Newey and Powell (2003): for all measurable  $g$ , we have:

$$E[g(t, x)|x, z] = 0 \ \forall (x, z) \in \text{supp}(X, Z) \Rightarrow g(t, x) = 0$$

where the equivalence comes from the fact that the LHS function is just  $T(g)$ .

# Consistency

Our is just a minimum distance estimator conditional upon draws from the 1<sup>st</sup> stage

Consistency has  $\|\bar{g} - \bar{g}^*\| \rightarrow 0$  after averaging over the instrument distribution

The relevant theory is in Newey and Powell (03) and Chen and Pouzo (12)

Promise to grow both nets at rates  $k(n) \rightarrow \infty$ ,  $k/n \rightarrow 0$ , then approximation theorems for neural nets (e.g. Hornik 91) imply density around the truth

For treatment (and 1<sup>st</sup> stage) continuous, also need # mixture components to grow

# Ads Application

Taken from Goldman and Rao (2014)

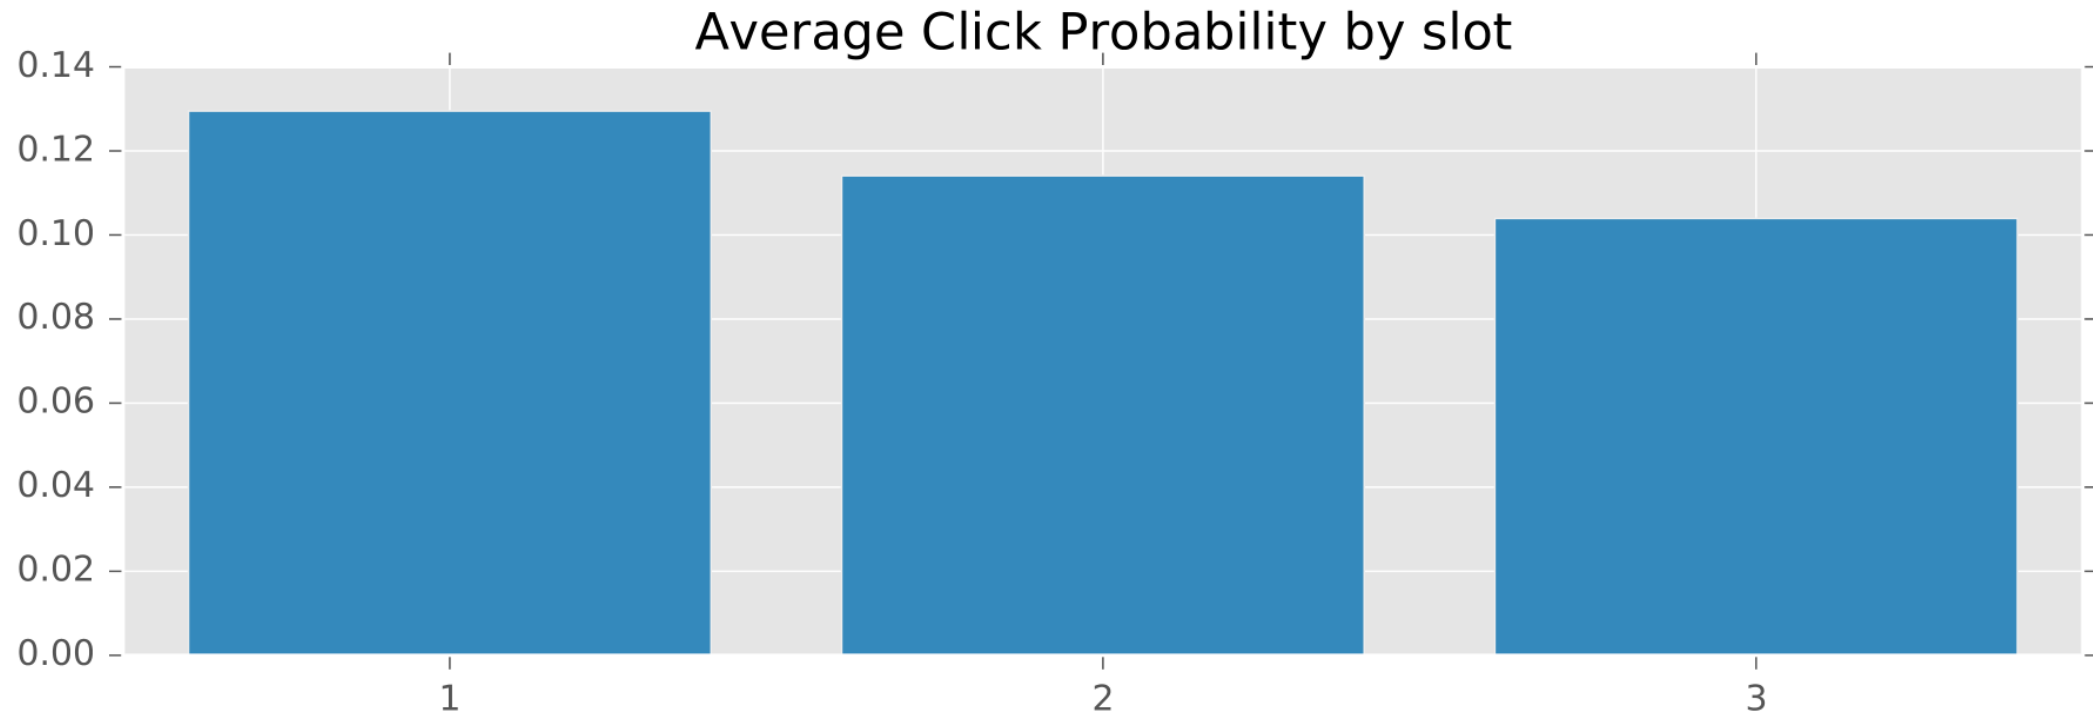
We have 74 mil click-rates over 4 hour increments for 10k search terms

Treatment: **ad position 1-3**

Instrument: **background AB testing (bench of ~ 100 tests)**


Covariates: **advertiser id and ad properties, search text, time period**

# Average Treatment Effects



These compare to observed click probabilities of 0.33, 0.1, and 0.05.


# Heterogeneous Treatment Effects



The *gaps* between ad-position lines can be interpreted *causally*.



# Heterogeneous Treatment Effects



The *gaps* between ad-position lines can be interpreted *causally*.

# Inference? Good question...

Recall the deep net

$$\hat{y}_i = \sum_k h_k^0(a_{ik}), \quad a_{ik} = \sum_j w_{kj} z_j, \quad z_j = \sum_l h_l^1(b_{il}), \dots$$

When training with SGD, we actually use **dropout** for regularization

At each update, calculate gradients against  $\tilde{\theta}$  where

$$\tilde{w}_{kj} = w_{kj} \delta_{kj}, \quad \delta_{kj} \sim \text{Bern}(\pi)$$

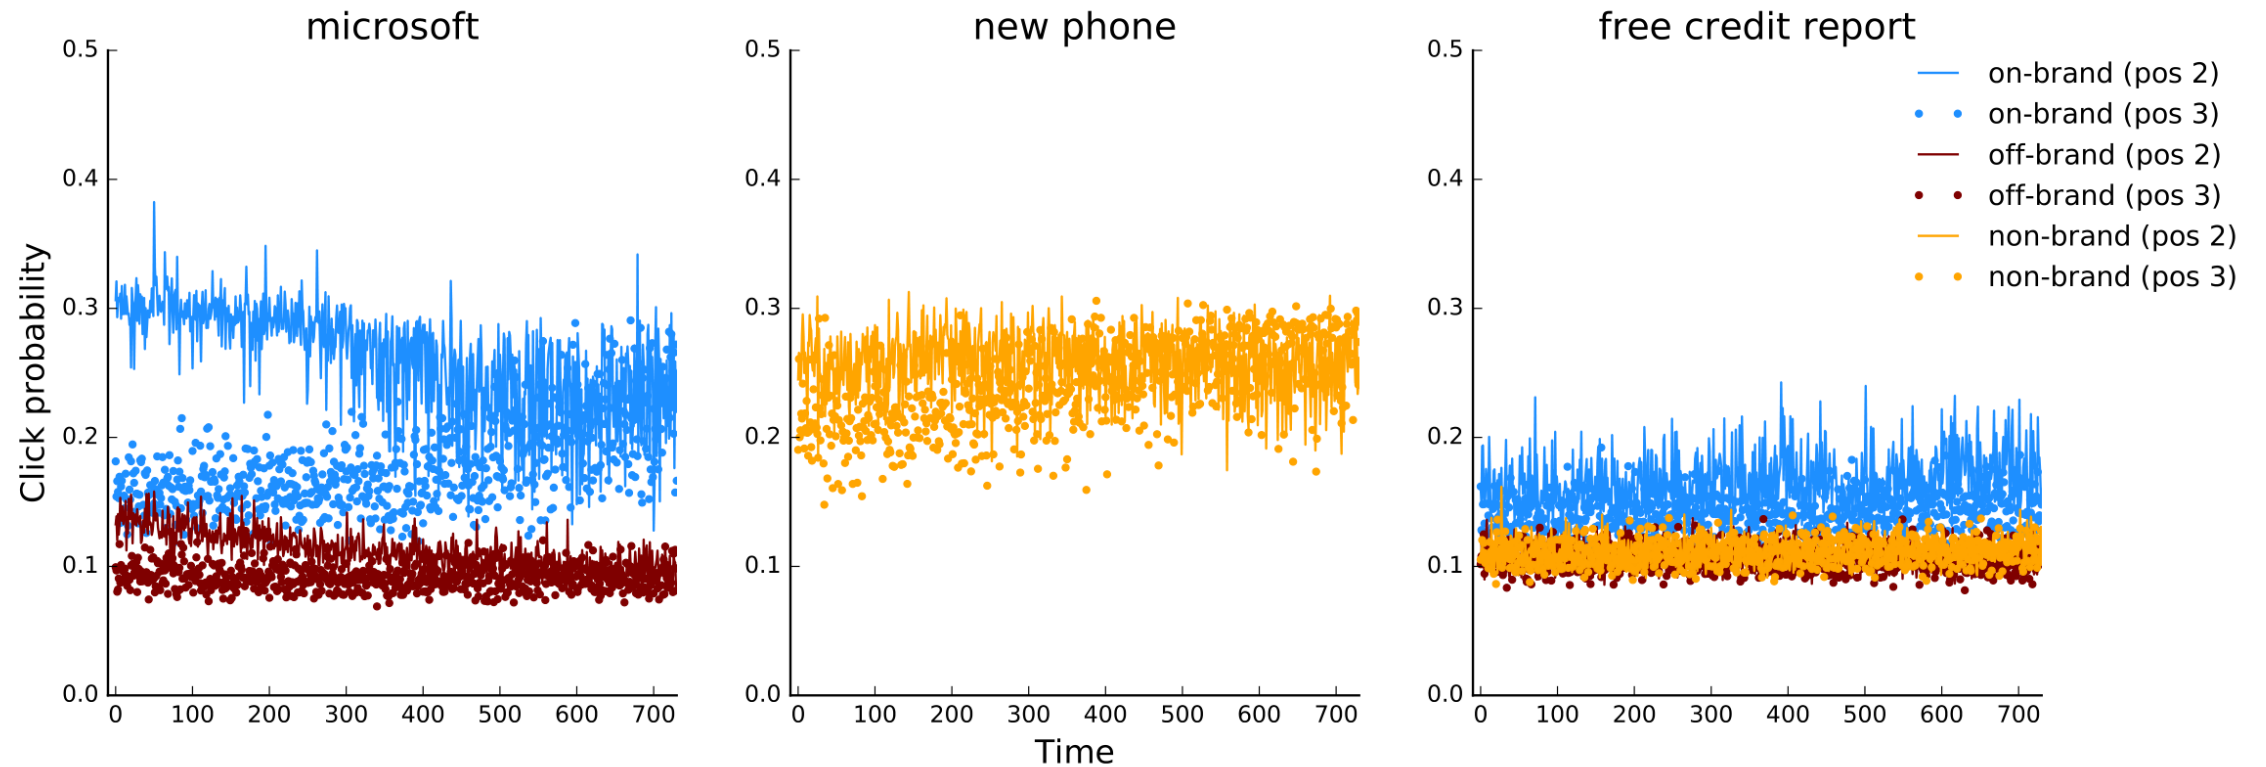
# Inference? Good question...

Variational Bayes minimizes

$$KL(q \mid p) = \mathbb{E}_q[\log q(y|x)] - \mathbb{E}[\log p(y|x, \mathbf{D})]$$

Gal and Ghahramani show that *dropout* SGD is optimizing variational distribution  $q$ , with uncertainty parametrized by the Bernoulli weights

# Posterior Treatment Effect Samples



Each point/dash is an independent draw from the `posterior`

# Economics and Artificial Intelligence

We have a track record pointing ML at questions of science + causation.

We're going to replicate this success *at scale on unstructured data*

We use economic theory to build systems of tasks that can be addressed with deep nets and other state-of-the-art ML.

This is the construction of systems for *Economic* Artificial Intelligence.