

Deep Counterfactuals via IV

Matt Taddy (Chicago + MSR) Greg Lewis (MSR)
Jason Hartford (UBC) Kevin Leyton-Brown (UBC)

Endogenous Errors

$$y = g(p, \mathbf{x}) + e \text{ and } \mathbb{E}[p e] \neq 0$$

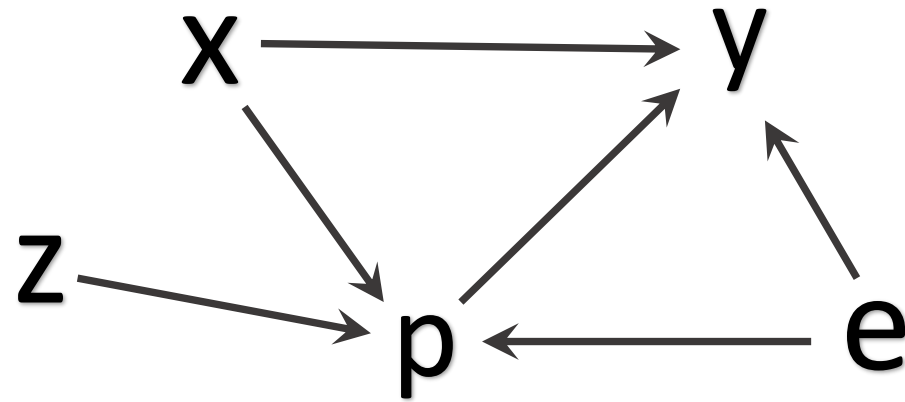
If you estimate this using naïve ML, you'll get

$$E[y|p, \mathbf{x}] = E_{e|p}[g(p, \mathbf{x}) + e] = g(p, \mathbf{x}) + E[e|p, \mathbf{x}]$$

This works for **prediction**. It doesn't work for **counterfactual** inference:

What happens if I change p independent of e ?


Instrumental Variables (IV)



In IV we have a special $z \perp e$ that influences policy p but not response y .

- Supplier costs that move price independent of demand (e.g., fish, oil)
- Any source of treatment randomization (intent to treat, AB tests, lottery)

Instrumental Variables (IV)



The *exclusion structure* implies

$$E[y|x, z] = E[g(p, x)|x, z] + E[e|x] = \int g(p, x) dF(p|x, z)$$

So to solve for *structural* $h(p, x)$ we have a new learning problem

$$\min_{g \in G} \sum \left(y_i - \int g(p, x_i) dF(p|x_i, z_i) \right)^2$$

$$\min_{g \in G} \sum \left(y_i - \int g(p, x_i) dP(p|x_i, z_i) \right)^2$$

2SLS:

$$p = \beta z + v \text{ and } g(p) = \tau p \text{ so that } \int g(p) dP(p|z) = \tau \hat{p} = \tau \hat{\beta} z$$

So you first regress p on z then regress y on \hat{p} to recover $\hat{\tau}$.

This requires strict assumptions and homogeneous treatment effects.

$$\min_{g \in G} \sum \left(y_i - \int g(p, x_i) dP(p|x_i, z_i) \right)^2$$

Or look to nonparametric 2SLS like in Newey and Powell:

$$g(p, x_i) \approx \sum_k \varphi_k(p, x_i) \text{ and } \varphi_k(p, x_i) \approx \sum_j \phi_{kj}(x_i, z_i)$$

But this requires careful crafting and will not scale with $\dim(x)$

$$\min_{g \in G} \sum \left(y_i - \int g(p, x_i) dF(p|x_i, z_i) \right)^2$$

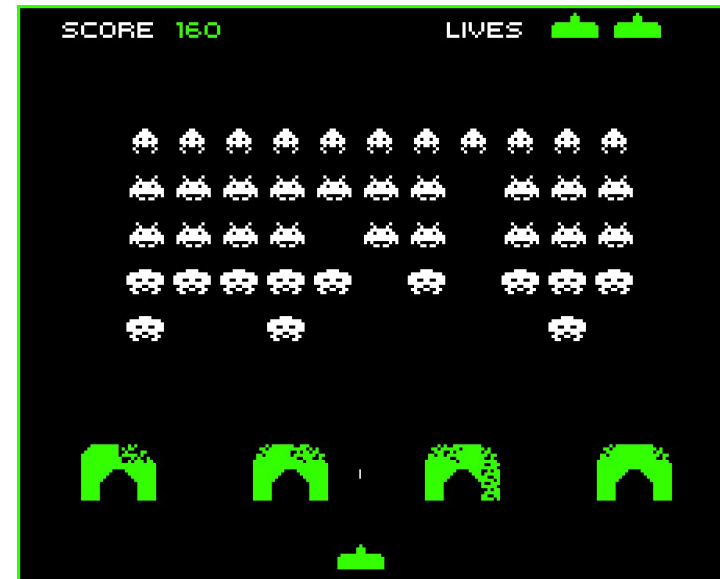
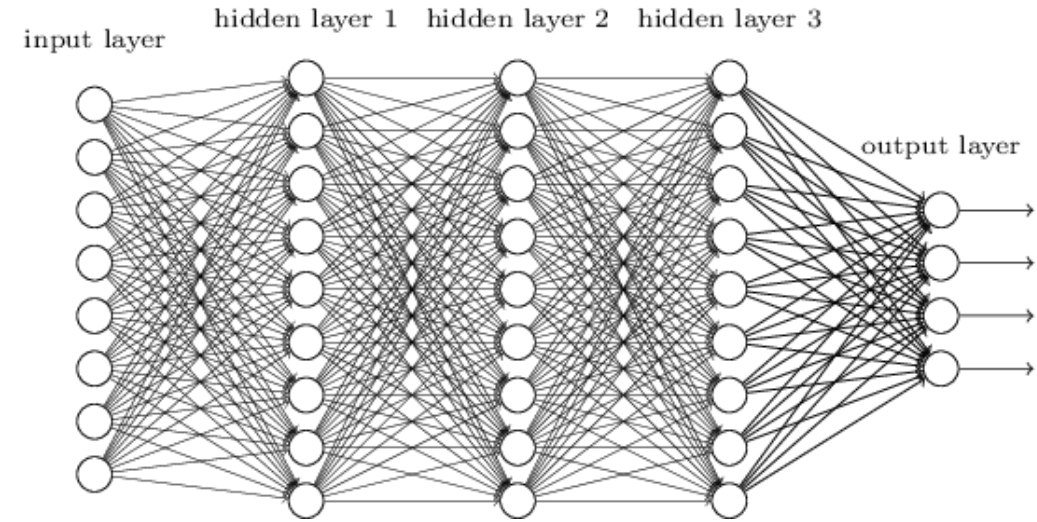
Instead, we propose to **target the integral loss function directly**

For discrete (or discretized) treatment

- Fit distributions $\hat{F}(p|x_i, z_i)$ with probability masses $\hat{f}(p_b|x_i, z_i)$
- Train \hat{g} to minimize $\left[y_i - \sum_b g(\hat{p}_b, x_i) \hat{f}(p_b|x_i, z_i) \right]^2$

And you've turned IV into two *generic* machine learning tasks

Learning to love Deep Nets



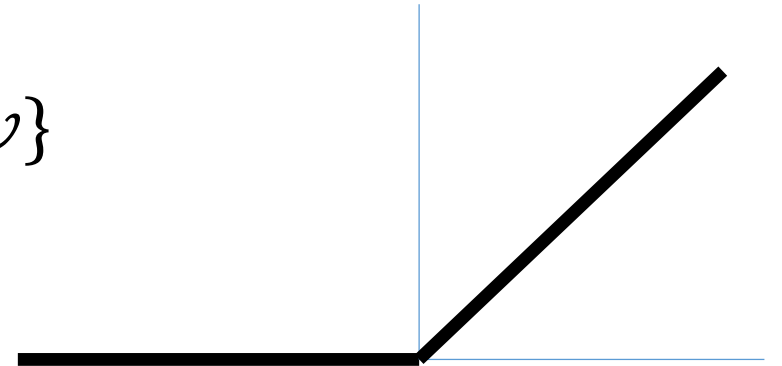
What is a deep net?

$$\hat{y}_i = \sum_k \eta_k(a_{ik}), \quad a_{ik} = \sum_j w_{kj} z_j, \quad z_j = \sum_l h_l^1(b_{il}), \dots$$

And so-on until you get down to a bottom layer $\{f_l(x_i)\}_l$



Many different variations here: recursive, convolutional, ...

Apart from the bottom, usually $h(v) = \max\{0, v\}$



e.g., first-stage learning for $F(p|x_i, z_i)$

Bishop 96: Final layer of network parametrizes a mixture of Gaussians



The second stage involves an integral loss function

If p is not discrete or can take many values, we can't just integrate

Brute force just samples from $\hat{F}(p|x_i, z_i)$ and you take gradients on

$$\frac{1}{N} \sum_i \left(y_i - \frac{1}{B} \sum_b g(\dot{p}_{ib}, x_i; \theta) \right)^2, \quad \dot{p}_{ib} \sim \hat{F}(p|x_i, z_i)$$

But this is super inefficient

Stochastic Gradient Descent

You have loss $L(\mathbf{D}, \theta)$ where $\mathbf{D} = [\mathbf{d}_1 \dots \mathbf{d}_N]$

In the usual GD, you iteratively descend

$$\theta_t = \theta_{t-1} - \mathbf{C}_t \nabla L(\mathbf{D}, \theta_{t-1})$$

In SGD, you instead follow *noisy* but *unbiased* sample gradients

$$\theta_t = \theta_{t-1} - \mathbf{C}_t \nabla L(\{\mathbf{d}_{t_b}\}_{b=1}^B, \theta_{t-1})$$

SGD for integral loss functions

Our one-observation stochastic gradient is

$$\nabla L(d_i, \theta) = -2 \left(y_i - \int g_\theta(p, x_i) d\hat{F}(p|x_i, z_i) \right) \int g'_\theta(p, x_i) d\hat{F}(p|x_i, z_i)$$

Do SGD by pairing each observation with *two independent* treatment draws

$$\nabla \hat{L}(d_i, \theta) = -2(y_i - g_\theta(\hat{p}_{i1}, x_i)) g'_\theta(\hat{p}_{i2}, x_i), \quad \hat{p}_{ib} \sim \hat{F}(p|x_i, z_i)$$

So long as the draws are independent, $\mathbb{E} \nabla \hat{L}(d_i, \theta) = \mathbb{E} \nabla L(d_i, \theta) = \nabla L(\mathbf{D}, \theta)$

Validation and model tuning

We can do *causal validation* via two OOS loss functions

Leave-out deviance on first stage

$$\sum_{i \in LO} -\log \hat{f}(p|x_i, z_i)$$

Leave-out loss on second stage (constrained fit on $\mathbb{E}[y|xz]$)

$$\sum_{i \in LO} (y_i - g_{\theta}(\dot{p}_i, x_i))^2, \quad \dot{p}_i \sim \hat{F}(p|x_i, z_i)$$


You want to minimize both of these (in order).

heterogeneous price effects


$$y = 100 + s\psi_t + (\psi_t - 2)p + e,$$

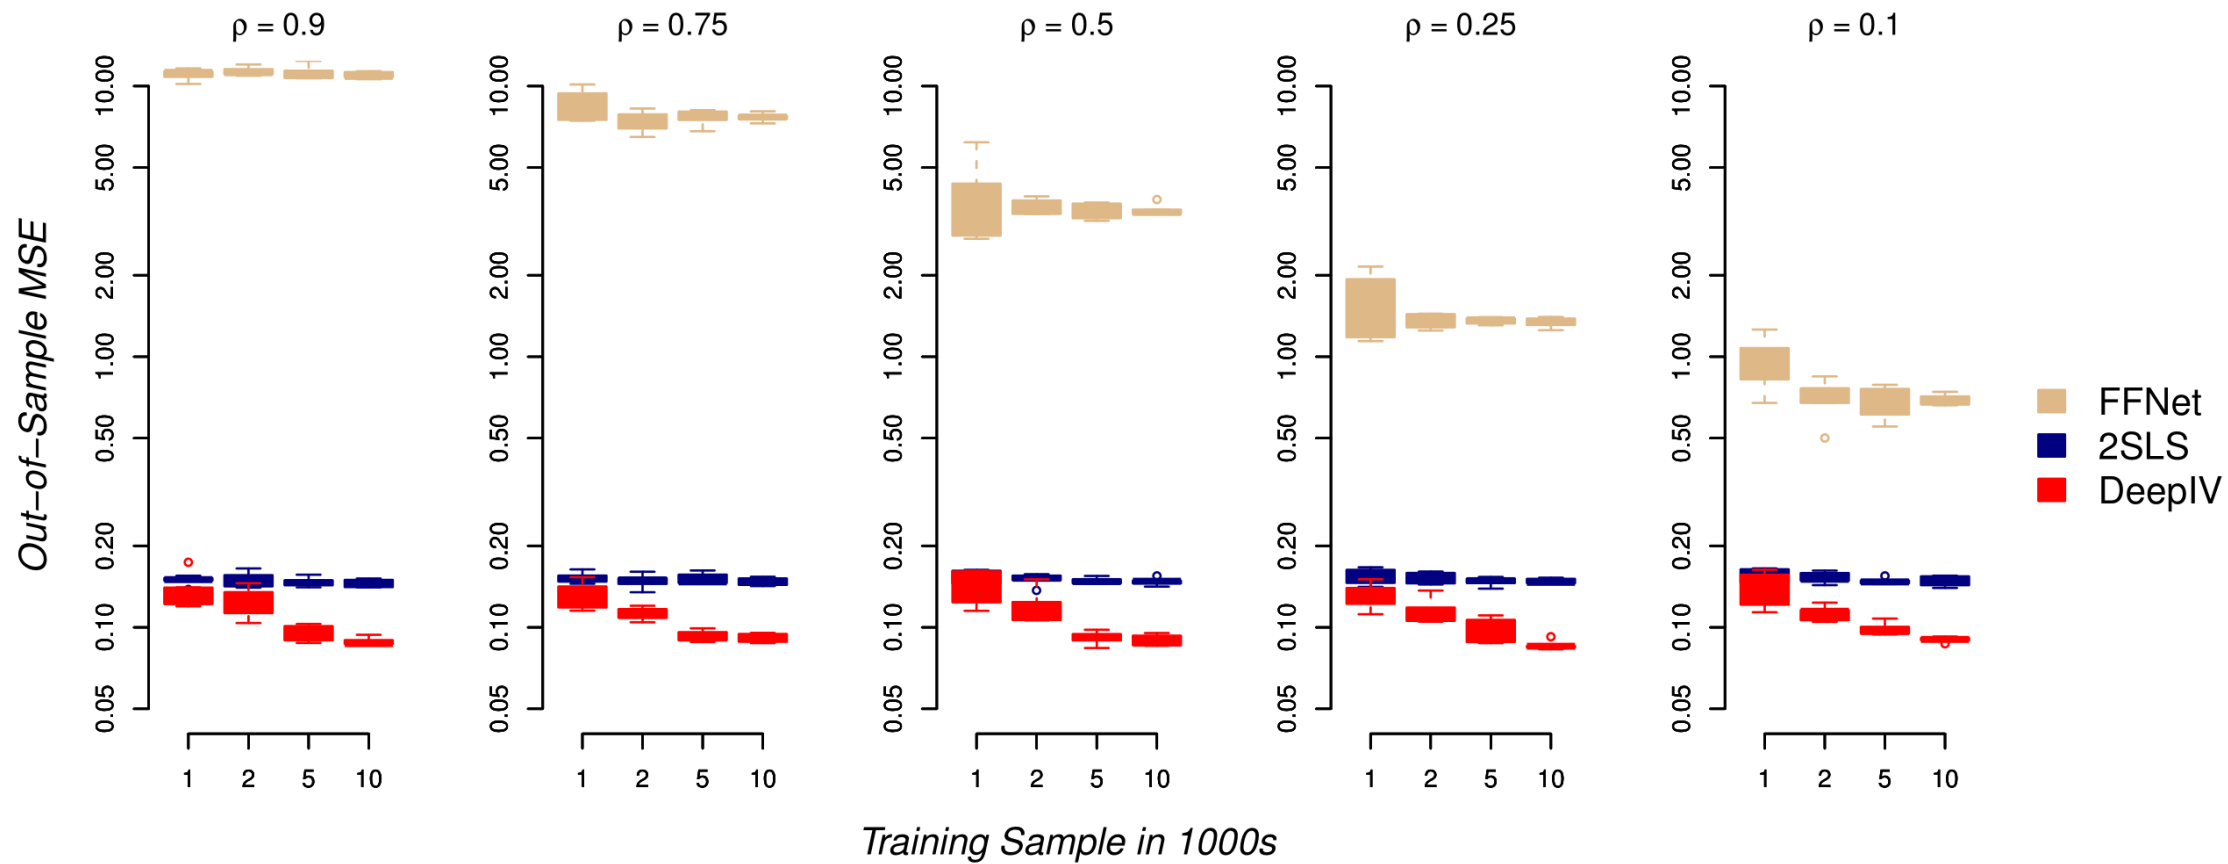
$$p = 25 + (z + 3)\psi_t + v$$

$$z, v \sim N(0, 1) \text{ and } e \sim N(\rho v, 1 - \rho^2),$$



- ‘time’ dependent prices and elasticity
- 7 different customer types multiply discounts, sensitivities, and demand





Inference? Good question

Data split! Get top node values and averages on left-out data:

$$\bar{\eta}_{ik} = E_{\hat{F}(\dot{p}|x_i,z_i)} \eta_k(x_i, \dot{p}) \quad \text{and} \quad \eta_{ik} = \eta_k(x_i, p_i)$$

Stack as instruments $\bar{H} = [\bar{\eta}_1 \cdots \bar{\eta}_L]'$ and treatments $H = [\eta_1 \cdots \eta_L]'$

Then the treatment effect is $\hat{\beta} = (\bar{H}'H)^{-1}\bar{H}'y$ with usual variance and

$$\text{var} \left(\hat{h}(x, p) \right) = \eta'(x, p) V_{\beta} \eta(x, p).$$

Inference? Good question

Or Bayes: recall the deep net

$$\hat{y}_i = \sum_k \eta_k(a_{ik}), \quad a_{ik} = \sum_j w_{kj} z_j, \quad z_j = \sum_l h_l^1(b_{il}), \dots$$

When training with SGD, we actually use **dropout** for regularization

At each update, calculate gradients against $W_l = \Xi_l \Omega_l$ at layer l where

$$\Xi_l = \text{diag}(\xi_{l1} \dots \xi_{lK_l}), \quad \xi_{kj} \sim \text{Bern}(c)$$

Variational Bayesian inference via dropout

VB minimizes $\mathbb{E}_q[-\log p(\mathbf{D}|\mathbf{W}) - \log p(\mathbf{W}) + \log q(\mathbf{W})]$

With $q(\mathbf{W}) = \prod_l \prod_k (c 1_{[W_{lk}=\Omega_{lk}]} + (1-c) 1_{[W_{lk}=0]})$ and normal prior, this objective becomes

$$\mathbb{E}_{q(c, \mathbf{\Omega})} \ell(D|W) + \sum_{l=1}^L c \lambda \|\mathbf{\Omega}_l\|^2 + \sum_{l=1}^L K_{l-1} [c \log(c) + (1-c) \log(1-c)].$$

So dropout is VB!

(more complex argument in Gal and Ghahramani 2015)

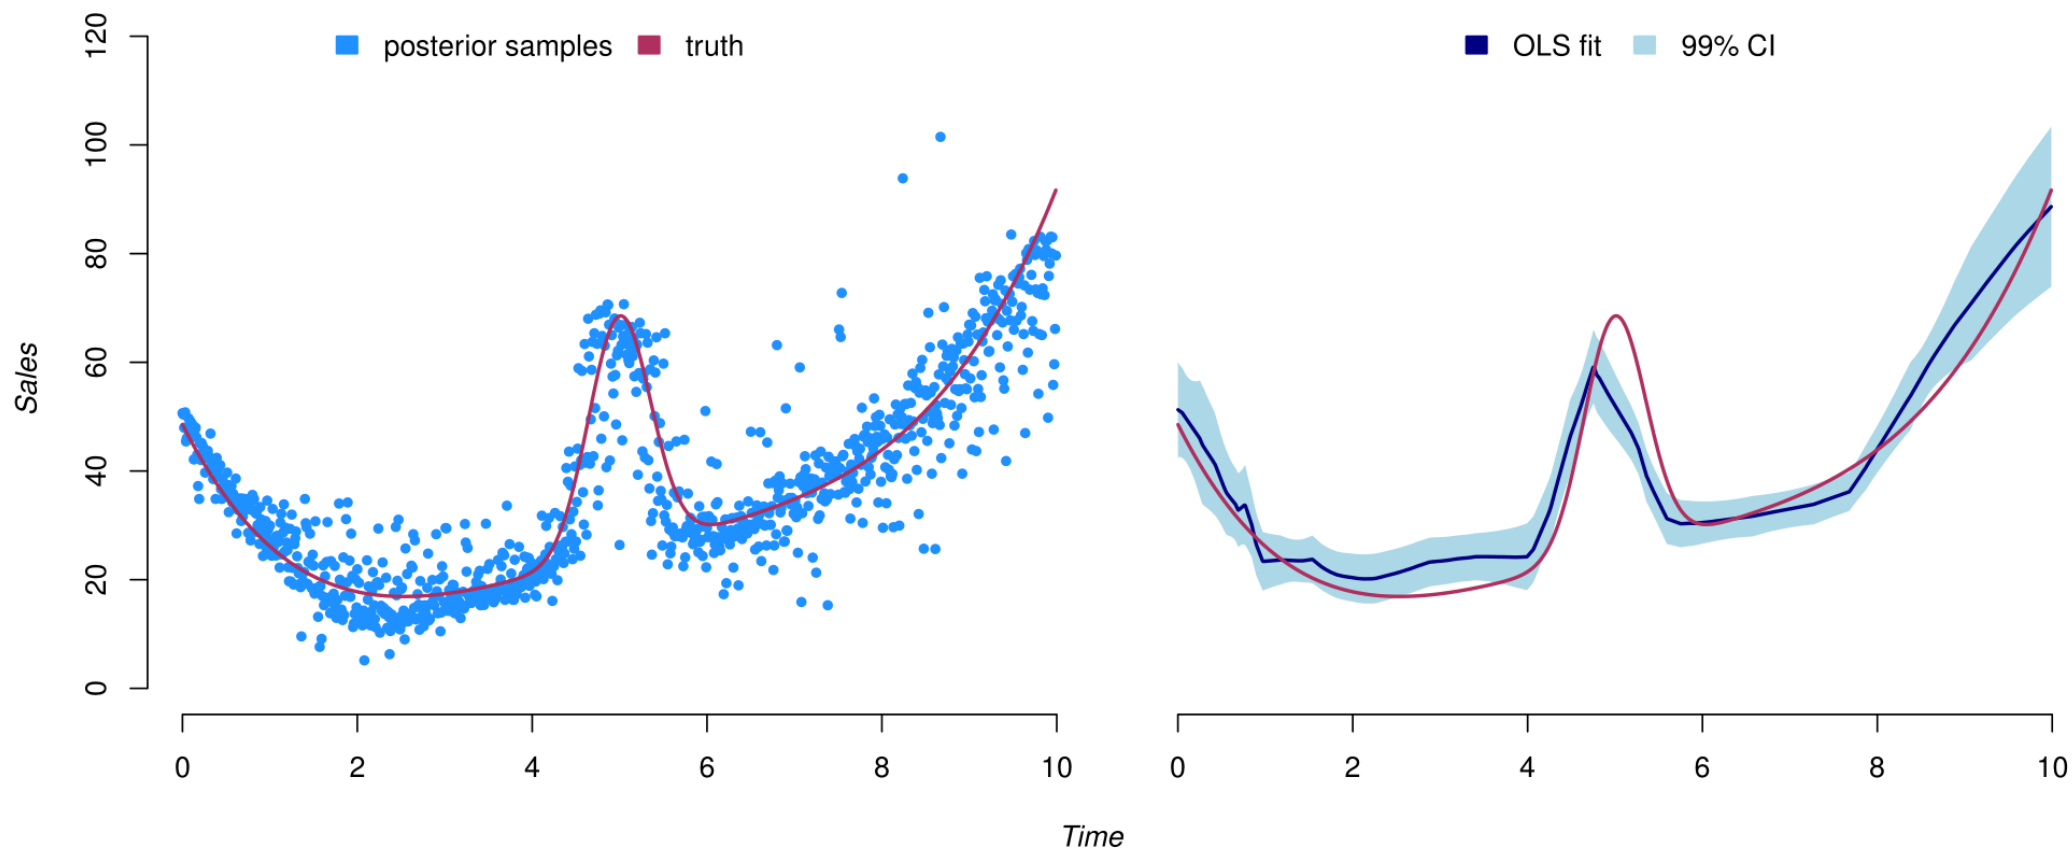
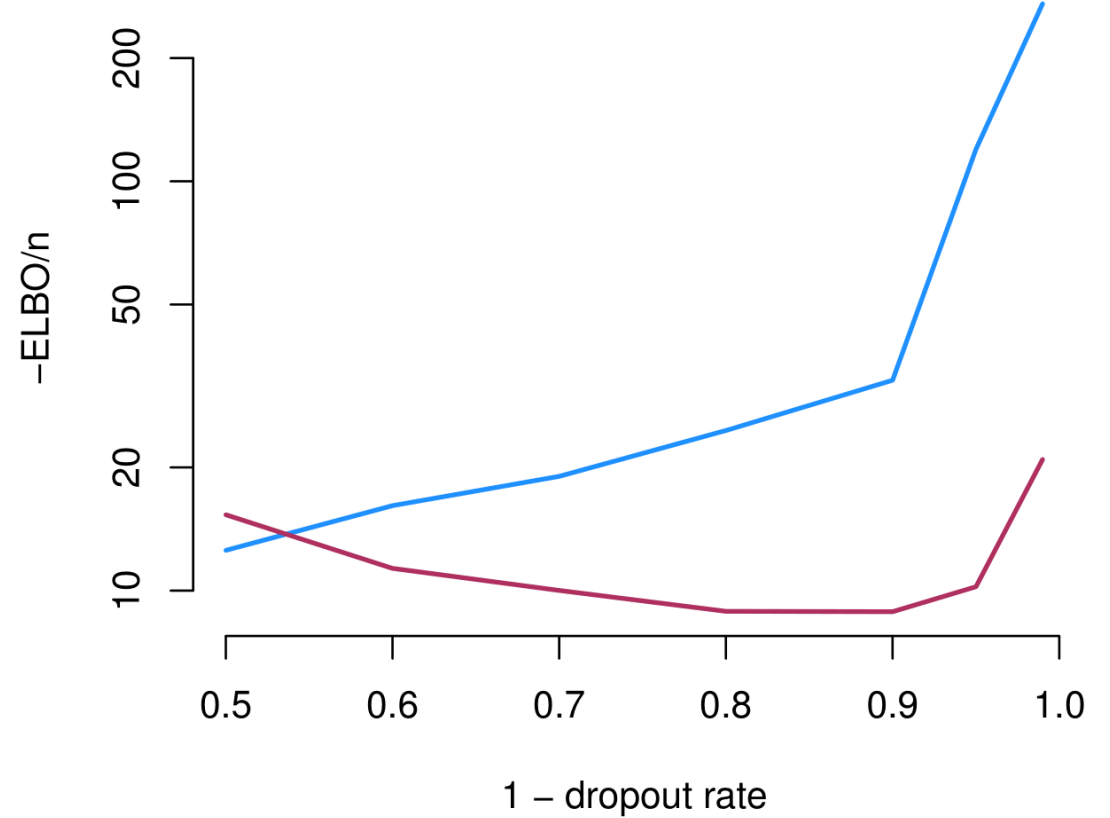
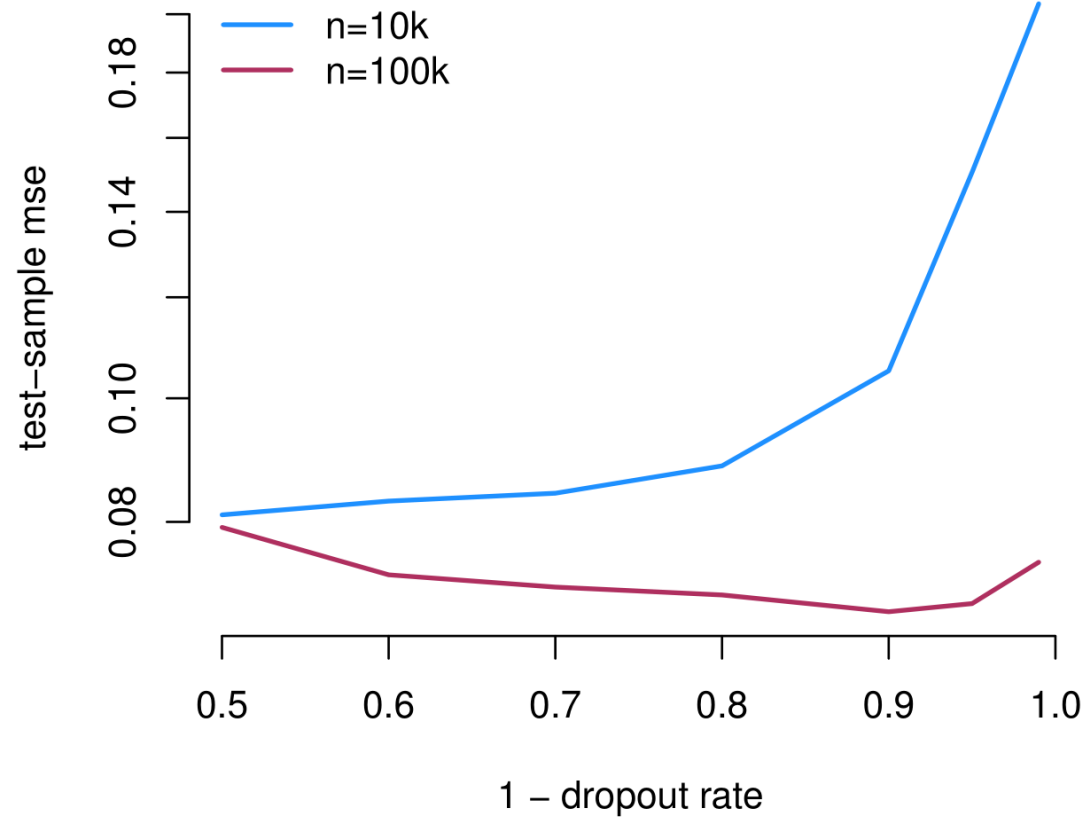


Figure 3: Bayesian (left) and Frequentist (right) inference for a central slice of the counterfactual function, taken at the average price and in our 4th customer category. Since the price effect for a given customer at a specific time is constant in (27), the curves here are a rescaling of the customer *price sensitivity* function.

Tuning the dropout rate is like treating it as a variational parameter



Ads Application

Taken from Goldman and Rao (2014)

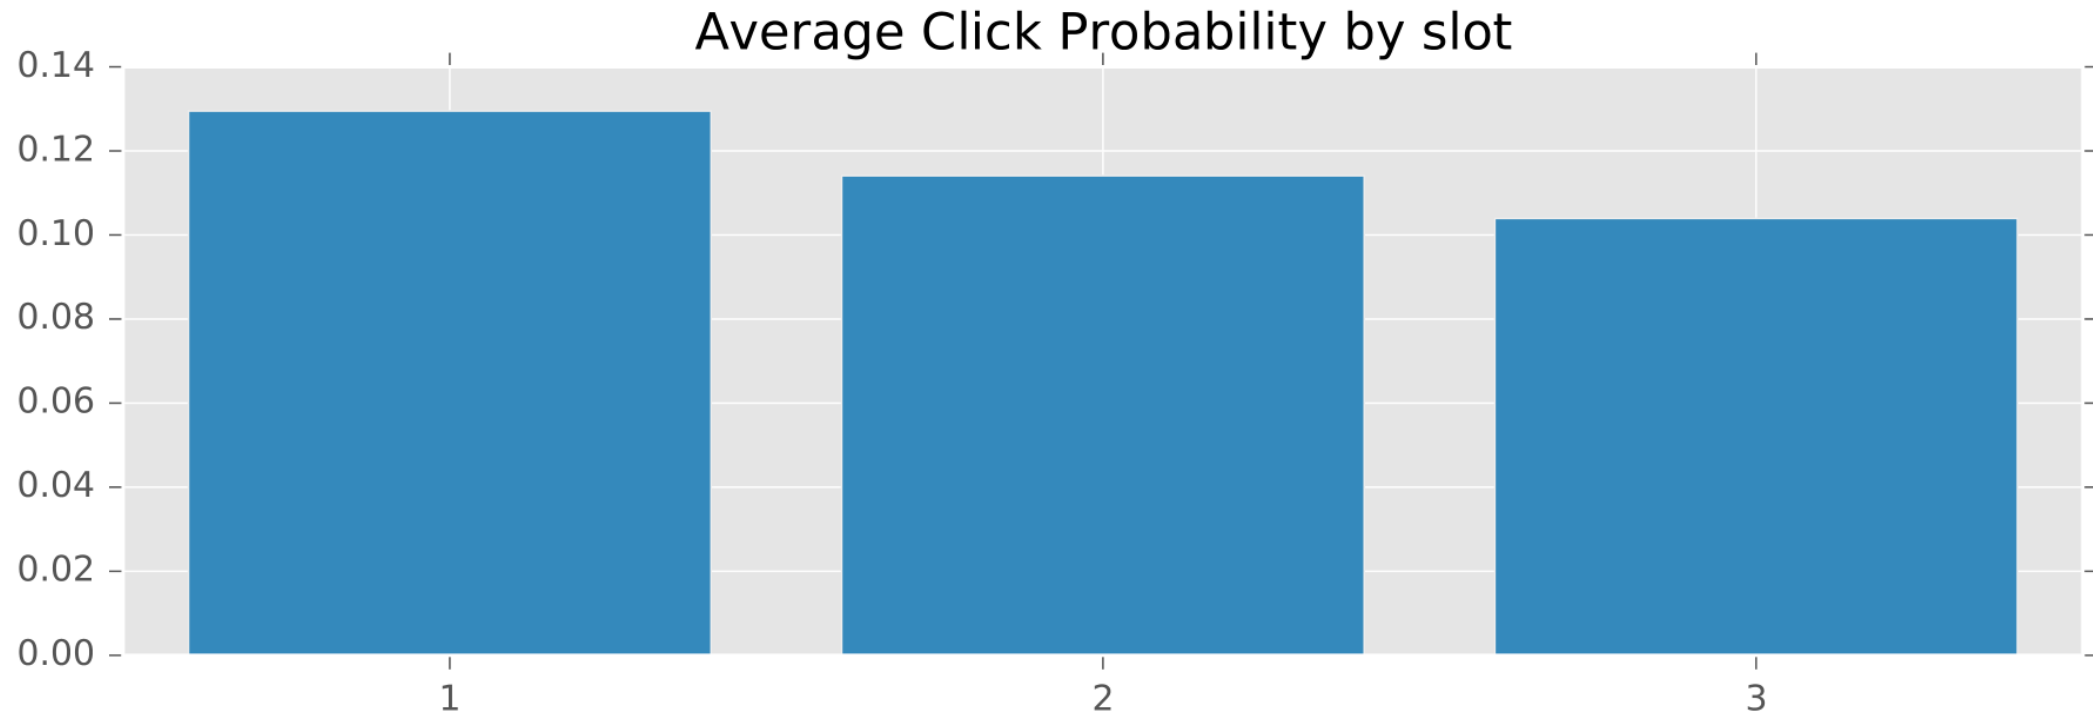
We have 74 mil click-rates over 4 hour increments for 10k search terms

Treatment: **ad position 1-3**

Instrument: **background AB testing (bench of ~ 100 tests)**


Covariates: **advertiser id and ad properties, search text, time period**

Average Treatment Effects



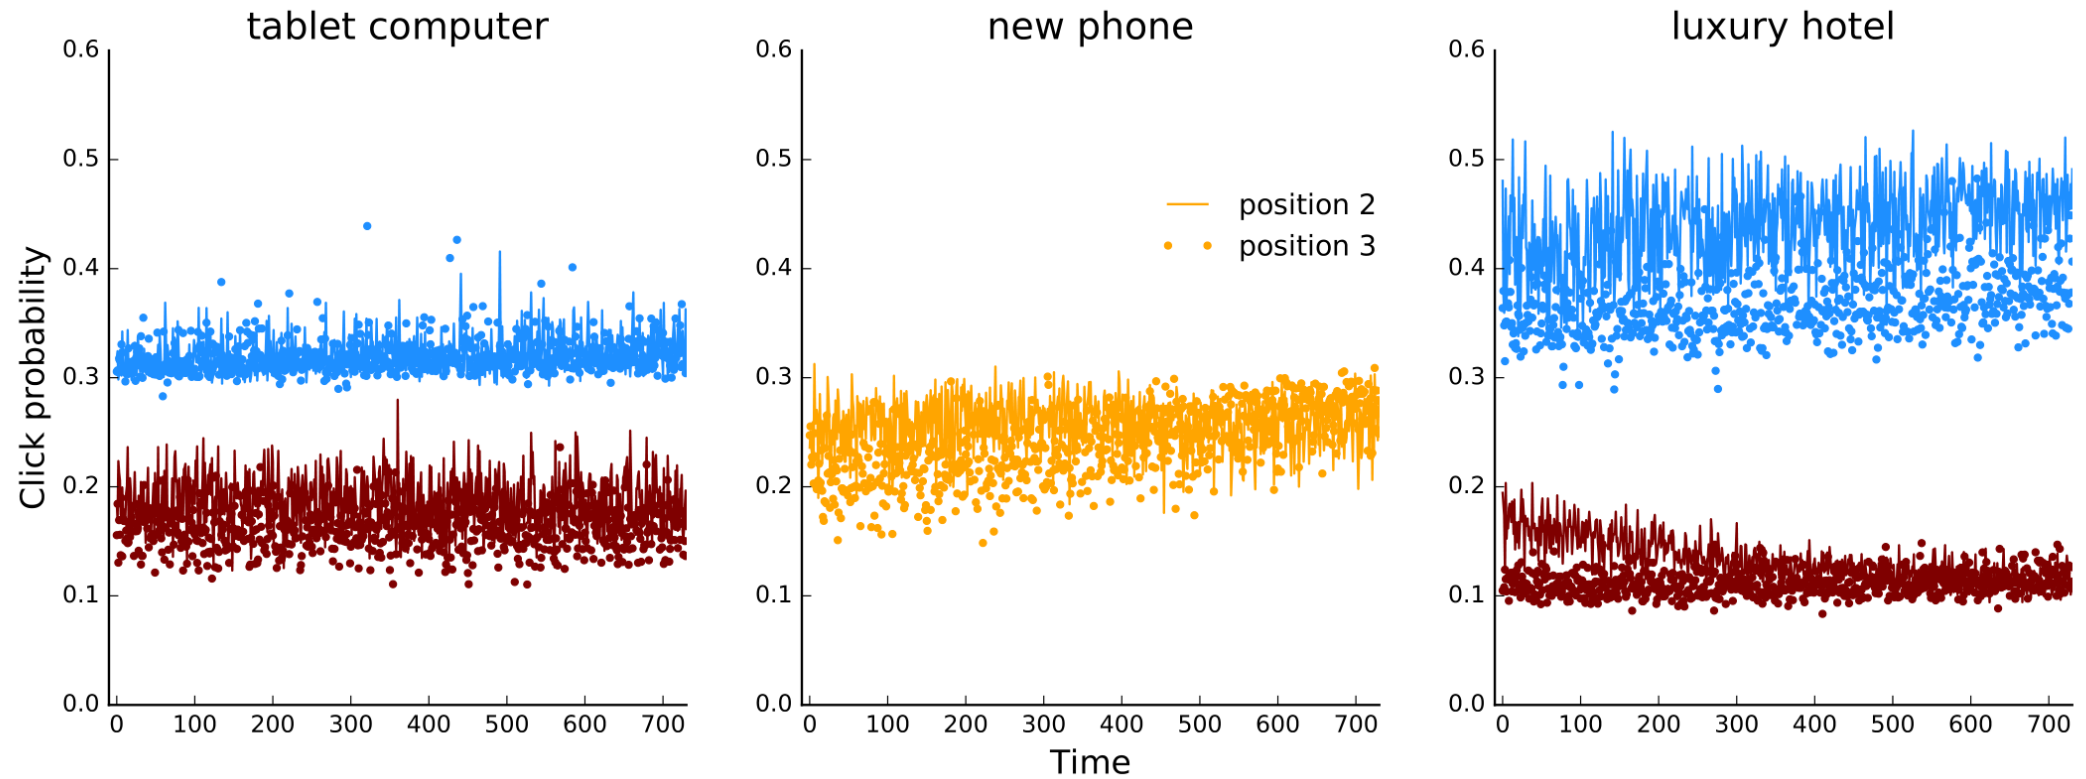
These compare to observed click probabilities of 0.33, 0.1, and 0.05.

Heterogeneous Treatment Effects



The *relative* click-probabilities can be interpreted *causally*.

Posterior Treatment Effect Samples



Each point/dash is an independent draw from the `posterior`

Alice

Established: December 5, 2016



Automated Learning and Intelligence for Causation and Economics

We use economic theory to build systems of tasks that can be addressed with deep nets and other state-of-the-art ML.

This is the construction of systems for *Economic AI*